

# Advanced Spring

## Distributed Tracing with Spring


**Boris Fresow, Markus Günther**

Adesso eduCamp 2023

Mastichari, Kos

## Some difficult questions

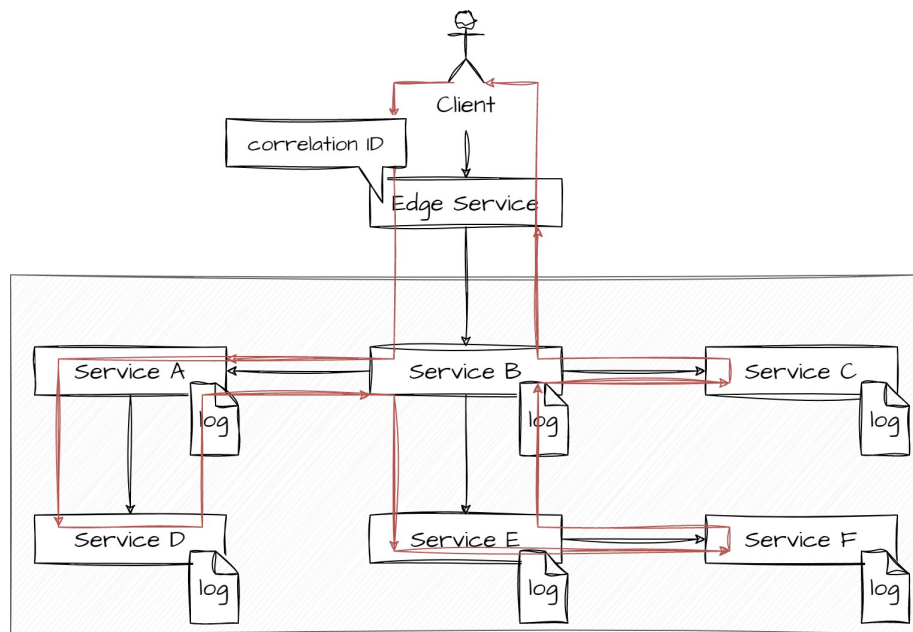
- Which services are involved in a use-case?
- Where are our bottlenecks?
- Why do we see so many 4xx / 5xx in our logs?
- Are we hitting our SLOs/NFRs?



*Distributed tracing provides end-to-end visibility into the performance and behavior of requests across interconnected services in a distributed system, enabling faster troubleshooting and optimization of the application.*

# The solution

- A unique ID that is passed from system to system
- All systems log into a central component
- Traces with the same ID are combined into a single trail



## Key Points

- (Enhanced) Observability
  - Comprehensive view of interaction between services
- Faster root cause analysis
  - Especially for Bottlenecks / Latency issues
- Service Dependency Visibility
  - Get a per use-case view of all involved services
- Anomaly Detection
  - Basis for monitoring, especially for non-functional requirements

# Components

# Conceptual Components

- **Instrumentation**

- Capture & record information about performance characteristics

- **Propagation**

- Provide the necessary information to correlate further operations
- Pass HTTP-Headers for HTTP Requests
- Log the span- / trace-id with the service logs for later analysis

- **Collection**

- Aggregate the collected data for a context

- Send it to a central component



## Conceptual Components (cont.)

- **Storage & Processing**
  - Persist data received
  - Correlate the related datasets
- **Visualization & Analysis**
  - User Interface for the data
  - Capability to search & filter
- **Alerting & Integration**
  - Automated analysis based on the collected data

- **Brave:** Tracer Library for Instrumentation (developed and maintained by Zipkin)
- **Spring Cloud Sleuth:** Layer on top of Brave for smooth Spring integration
  - *No longer available for Spring Cloud 2022.x*
- **Micrometer Observation:** Vendor/technology agnostic metrics collector for Java
  - Can be used with Brave & Zipkin or other technologies
  - *The new default for Spring Cloud 2022.x*
- **Zipkin:** Distributed Tracing System for collection & visualization

- **Trace:** Unique object that contains 1..n spans - builds a latency tree
- **Span:** Single host-view of an operation. Most important fields are:
  - `traceId`: the root trace ID
  - `parentId`: the ID of the parent span (null if root)
  - `name`: the name of the span
- **Tag:** A marker for a span - can be used in a query
- **Annotation:** A specific point in time in a span with an attached string value

# Spring Cloud Sleuth

*Spring Cloud Sleuth will not work with Spring Boot 3.x onward.  
The last major version of Spring Boot that Sleuth will support is  
2.x.*

We'll only take a quick look at Sleuth for the sake of legacy knowledge

- Easy integration via Spring Cloud dependencies

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-sleuth</artifactId>
</dependency>
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-sleuth-zipkin</artifactId>
</dependency>
```

- This includes the auto configuration part that ...
  - exposes Java beans for customizing data collection & presentation
  - decorates existing beans (e.g. `RestTemplate` for propagation)

# Configuration

- The default configuration should be changed - for our example we used this config

```
spring.application.name: Client Application

spring:
  zipkin:
    base-url: ${ZIPKIN_API_ENDPOINT:http://127.0.0.1:9411/}
  sleuth:
    sampler:
      probability: 1.0
```

- Application name will be reported with the spans
- The Base-URL is where the traces will be sent to
- The probability default is 0.1 - so only 10% of traces would be gathered
  - Crucial overload protection for high load systems

## Creating Traces

- Spans will be created and reported implicitly for supported operations
  - `@Scheduled`, `@RestController`, ...
  - The current span can also be modified
- Can also be done explicitly
  - Via annotations
  - Programmatically



## Creating Traces via Annotations

- `@NewSpan("name")` tells Sleuth to open a new span with the given name
  - Will be the child of the already existing span - if there is one
- `@ContinueSpan` can be used to annotate the current span
- `@SpanTag` can be used on parameters to add them as tags

```
@NewSpan("user-lookup")
public User lookupUserByName(@SpanTag(key = "username") final String username) {
    (...)
}
```

## Creating Traces programmatically

- `@NewSpan("name")` tells Sleuth to open a new span with the given name
  - Will be the child of the already existing span - if there is one
- `@ContinueSpan` can be used to annotate the current span
- `@SpanTag` can be used on parameters to add them as tags

```
Span newSpan = null;
try (Tracer.SpanInScope ws = this.tracer.withSpanInScope(initialSpan)) {
    newSpan = this.tracer.nextSpan().name("calculateCommission");
    newSpan.tag("username", username)
} finally {
    if (newSpan != null) {
        newSpan.finish();
    }
}
```

# Propagation

- Sleuth will add the tracing context to outgoing requests
  - will work out of the box for directly supported mechanisms (e.g. RestTemplate)
- For HTTP requests HTTP-Headers (x-b3-) will be added
- Instrumentation library on the receiving side will continue the context

```
> p authentication = {UsernamePasswordAuthenticationToken@8562}
v p headers = {LinkedMultiValueMap@8563} size = 9
  > "accept" -> {ArrayList@8580} size = 1
  > "authorization" -> {ArrayList@8582} size = 1
  > "x-b3-traceid" -> {ArrayList@8584} size = 1
  > "x-b3-spanid" -> {ArrayList@8586} size = 1
  > "x-b3-parentspanid" -> {ArrayList@8588} size = 1
  > "x-b3-sampled" -> {ArrayList@8590} size = 1
  > "user-agent" -> {ArrayList@8592} size = 1
  > "host" -> {ArrayList@8594} size = 1
  > "connection" -> {ArrayList@8596} size = 1
```

# Micrometer Observation

# Micrometer Observation

- Micrometer is a very powerful library for collecting metrics
  - You will hear more about this tomorrow!
- The Micrometer Observation API is vendor neutral
  - Modular approach allows a *do once, use many times* approach.
  - Has a bridge to Brave and therefore Zipkin

# Integration

- Dependencies are not from Spring
- Requires out-of-band dependency management
- Some integration with spring-boot-starter-actuator for autowiring beans

```
<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-observation</artifactId>
</dependency>
<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-tracing-bridge-brave</artifactId>
</dependency>
<dependency>
  <groupId>io.zipkin.reporter2</groupId>
  <artifactId>zipkin-reporter-brave</artifactId>
</dependency>
```

## Integration (cont.)

- Has a dedicated dependency for test support
  - Very useful - we will take a look at this tomorrow

```
<dependency>  
  <groupId>io.micrometer</groupId>  
  <artifactId>micrometer-observation-test</artifactId>  
  <scope>test</scope>  
</dependency>
```

- `Observation` is a wrapper class for an observable operation
- Observations follow a life cycle
  - `start`: The operation has begun execution
  - `scope started`: A scope within the operation has been opened
  - `scope ended`: A scope was closed
  - `event`: An event has happened while observing (e.g. annotations in Sleuth)
  - `error`: The operation threw an error
  - `stop`: The operation has ended



- Observations are handled by `ObservationHandlers`
  - The signature is `ObservationHandler<T extends Observation.Context>`
  - The interface provides handling methods for all lifecycle events
- Micrometer provides an `ObservationRegistry` to register `ObservationHandler`
  - `ObservationHandler` are called when they support a specific `Observation.Context`
- The `Observation.Context` is a mutable data holder to share data between states
  - The context will be propagated between threads!

# Configuration

```
spring.application.name: Micrometer Tracing
logging.pattern.level: "%5p [%${spring.zipkin.service.name:${spring.application.name}}],%X{traceId:-},%X{spanId:-}"

management:
  tracing:
    propagation:
      type: b3
    sampling:
      probability: 1.0
  zipkin:
    tracing:
      endpoint: ${ZIPKIN_API_ENDPOINT:http://127.0.0.1:9411/api/v2/spans}
```

- Probability and endpoint are similar (plus path in URL) to Sleuth
- We need to reconfigure the logging pattern manually to log trace- and span-IDs
- Propagation type b3 for interop with Sleuth (default is w3c)

## Configuration (cont.)

- To use the `@Observed` annotation we need to register the Bean for AOP
- The `ObservationRegistry` is the registry responsible for observation state management

```
@Bean
public ObservedAspect observedAspect(ObservationRegistry observationRegistry) {
    return new ObservedAspect(observationRegistry);
}
```

# Creating Observations via Annotations

- `@Observed` tells Micrometer that the method should be observed

```
@Observed(  
    name = "time.reporting",  
    contextualName = "time-report",  
    lowCardinalityKeyValues = {"class.name", "TimeReportingTask"}  
)  
public void reportTimeTask() {  
    (...)  
}
```

## Creating Observations programmatically

- `Observation` offers several static factory methods to create a new `Observation`
- `observe` will
  - start/stop the `Observation` (if not already done)
  - Open/close a new scope
  - Catch all errors

```
void observeWork(@Autowired ObservationRegistry registry) {  
    final var context = new Observation.Context().put(String.class, "information");  
    // The context is optional  
    Observation  
        .createNotStarted("operation.name", () -> context, registry)  
        .observe(this::doWork);  
}
```

## Manipulate Observations

- The `ObservationRegistry` is the global holder for all `Observations` (comparable to the Tracer in Sleuth)
- Fetch current observation via `registry.getCurrentObservation()`
  - Will be `null` if we're not being observed

```
void doWork(@Autowired ObservationRegistry registry) {  
    final var observation = registry.getCurrentObservation();  
    (...)  
}
```

## Observations Bridge

- An `Observation` conceptually has nothing to do with a trace
- `micrometer-tracing` provides a set of `ObservationHandler` to handle `Observations`
  - Transitive dependency of `micrometer-tracing-bridge-brave`
- `micrometer-tracing-bridge-brave` is responsible for
  - converting the traces to Brave
  - Reporting to Zipkin via Brave

# Zipkin



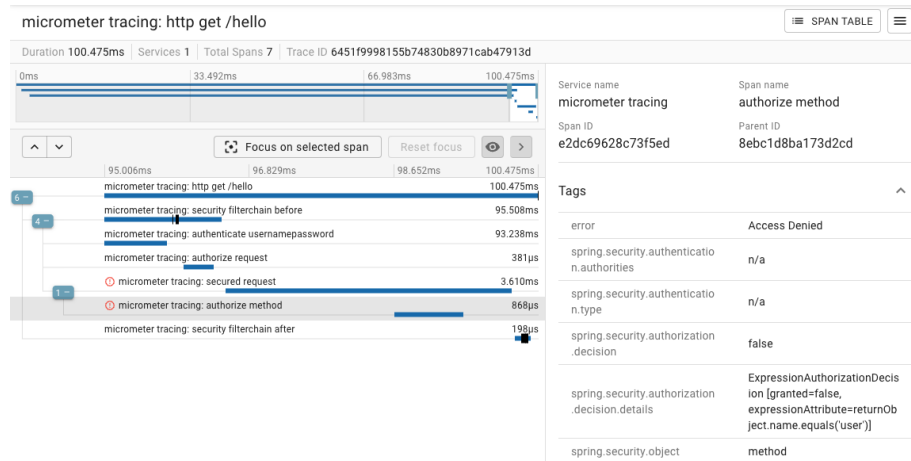
*Zipkin is a distributed tracing system. It helps gather timing data needed to troubleshoot latency problems in service architectures. Features include both the collection and lookup of this data.*

## Zipkin (cont.)

- Zipkin provides
  - a HTTP Facade to receive (and query!) spans
  - an UI for visualization & querying
- Is usually run as a standalone application
  - but can be integrated in a service as well
- Spans can be collected via HTTP, Kafka or RabbitMQ
- Storage options are in-memory, MySQL, Cassandra or Elasticsearch
  - There are a lot of community supported plugins e.g. for SQS, Kinesis

## Lab assignments will focus on

- visualization & querying of data
- tracing HTTP Requests



**Questions?**

**Lab**

It's time to get started!

Let's take a look at the repository and README.md



